

Session 2: Kokkos


Outline

- Computing architectures, shared- vs. distributed-memory parallelism
- Initialization
- Views
- Parallelizing for loops and reductions
- Random number generation
- ScatterViews
- Virtual functions

Parallelism comes in two flavors


Shared-memory

- Participating **threads** all share a chunk of memory
- Data (variables) accessible to all
- Impossible to scale up to supercomputer size



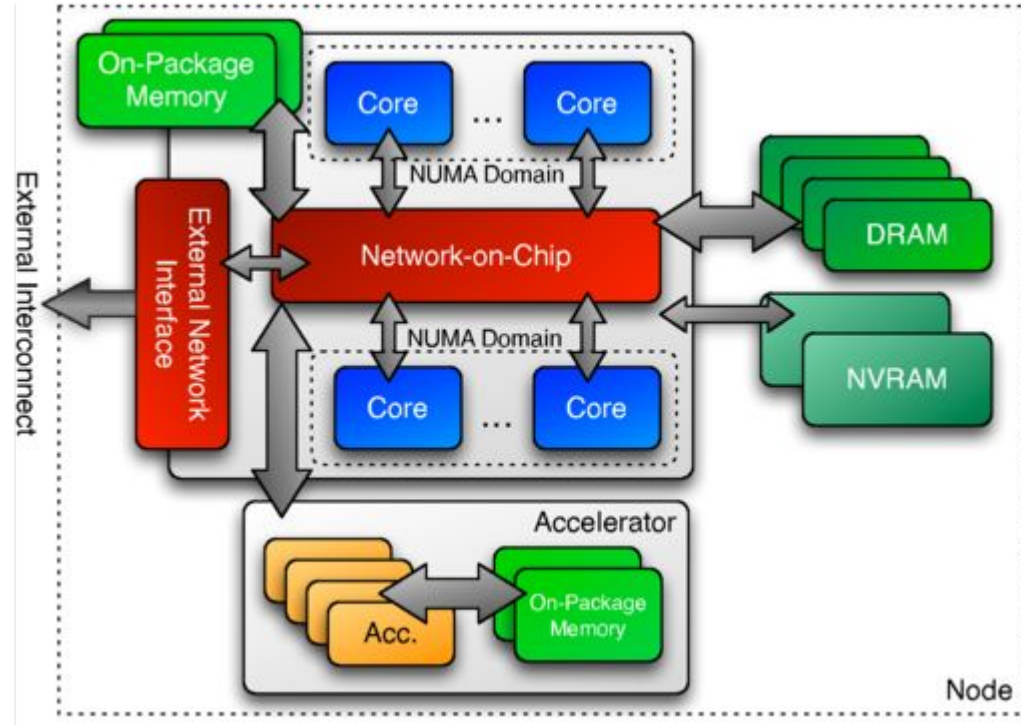
Distributed-memory

- Participating **tasks** or **processes** each have their own chunk of memory
- Data must be manually communicated between tasks



Modern supercomputers require us to use both

- Supercomputers comprise many nodes, which are independent computers in their own right
- Each node has its own memory
- Nodes may contain many processors and even GPUs, which have their own separate memory




Kokkos lets us exploit shared-memory parallelism on any system

- GPU manufacturers each have their own little languages that they prefer
 - NVIDIA uses CUDA
 - AMD uses HIP
 - Intel uses DPC++
- Kokkos basically translates code to each of these, depending on how you install Kokkos
- Rather than writing a bunch of special cases, we can write one code and trust that Kokkos optimizes it for the target architecture (and it does this well)
- There are many alternatives (OpenCL, OpenGL, SYCL, etc.), but we use Kokkos because it *emphasizes scientific computing* on typical *supercomputers in the US*
- Kokkos supports arbitrarily many CPUs on one memory bank and up to one GPU at a time
 - We'll need MPI to get more than one device (next week!)

The host/device model is used by Kokkos

- The set of CPUs available to the program and attached to main RAM are called the “host”
- Attached accelerator (GPU, FPGA) is called the “device”
- Device memory is separate from host memory and must be manually communicated
- Device memory accessible by all device threads, host memory accessible by all host threads



Initialize Kokkos in code before doing anything else

- ... and finalize before ending anything!
- Kokkos will request resources from the system, like GPU access
- By feeding in argc and argv, Kokkos will also “trim out” its own command-line options, leaving just ones specific to your program
 - To see the options, run any program built with Kokkos with the `-kokkos-help` option
- Can also use the ScopeGuard, which is just a class with initialize in its ctor and finalize in its destructor
 - Obviates having to finalize manually, since it's called when ScopeGuard goes out of scope

```
#include <Kokkos_Core.hpp>

int main(int argc, char *argv[]) {
    Kokkos::initialize(argc, argv);
    {
        // your code here
    }
    Kokkos::finalize();
    return 0;
}

int main(int argc, char *argv[]) {
    Kokkos::ScopeGuard scope_guard(argc, argv);
    {
        // your code here
        // no need to finalize!
    }
    return 0;
}
```

Why do we put our code in curly braces?

- Kokkos::finalize must be called only after ALL Kokkos data has been deallocated!
- Without curly brace scope guards, data will be destroyed at program termination, AFTER Kokkos::finalize
- The Kokkos::ScopeGuard helps here too
 - If it's constructed first, it'll be destroyed last, ensuring that Kokkos::finalize is called later

```
int main(int argc, char *argv[]) {
    Kokkos::initialize(argc, argv);
    {
        Kokkos::View<double> example;
    }
    // example is now destroyed, all is good
    Kokkos::finalize();
    return 0;
}

int main(int argc, char *argv[]) {
    Kokkos::initialize(argc, argv);
    Kokkos::View<double> example;
    Kokkos::finalize();
    return 0;
    // example isn't destroyed until now. Bad!
}

int main(int argc, char *argv[]) {
    Kokkos::ScopeGuard scope_guard(argc, argv);
    Kokkos::View<double> example;
    return 0;
    // example isn't destroyed until now,
    // but scope_guard is destroyed even later, so all good!
}
```


Kokkos defines “execution spaces” at compile time

- These specify places where work can happen and data can be stored
- Execution spaces are enabled in Kokkos as options when Kokkos is installed
- Kokkos also defines two default execution spaces
 - `DefaultExecutionSpace` is set to the highest space in the hierarchy to the right
 - `DefaultHostExecutionSpace` is set to the highest host space to the right
 - Note that, if you do not enable devices in Kokkos, these can be the same!
- Default spaces enable code portability; never use specific spaces!

Hierarchy of execution spaces

1.	<code>Kokkos::Cuda</code> (NVIDIA)	Devices
2.	<code>Kokkos::Experimental::OpenMPTarget</code> (generic GPUs)	
3.	<code>Kokkos::Experimental::HIP</code> (AMD)	
4.	<code>Kokkos::Experimental::SYCL</code> (Intel)	
5.	<code>Kokkos::OpenMP</code>	Host
6.	<code>Kokkos::Threads</code>	
7.	<code>Kokkos::Experimental::HPX</code>	
8.	<code>Kokkos::Serial</code>	

A View is a flexible, multidimensional data array

- Layout of data is determined at compile time for optimal performance in parallel
- Label string can be provided to provide more helpful error messages
- Dimensions of array set in datatype template argument
 - Number of asterisks sets number of dimensions known at runtime
 - Numbers in square brackets set dimensions at compile-time
 - Runtime dimensions must precede compile-time dimensions
- Kokkos::deep_copy can be used to copy between Views with the same dimensions and layout, or to copy a single value to all entries of a View

The first arg for the ctor for Views is a label string, and the remaining args are sizes for the runtime-sized dimensions

```
// Creating a 4x5x3 array with two runtime-sized dimensions
// and one compile-time-sized dimension.
Kokkos::View<double**[3]> example("label", 4, 5);

Kokkos::View<double**[3]> to_copy("to_copy", 4, 5);
// This copies the entries of to_copy to example.
Kokkos::deep_copy(example, to_copy);

// Set all the entries of example to 4.0
Kokkos::deep_copy(example, 4.0);
```

Views are not generally accessible from host

- So we can't just read data from a View
- This is because, by default, the data lives in the DefaultExecutionSpace, which may be on device
- But we can create a host mirror
- Behavior of host mirror differs based on enabled execution spaces
 - If DefaultExecutionSpace is a device space, host mirror is allocated on host and is separate memory from original View
 - If DefaultExecutionSpace is a host space, host mirror is a shallow copy to same data to avoid redundancy
- View data can only be safely accessed on host through host mirrors

Note that the host mirror datatype is just the View you want to mirror followed by `::HostMirror`

```
Kokkos::View<double*> example("example", 6);
Kokkos::View<double*>::HostMirror h_example =
    Kokkos::create_mirror_view(example);
// For performance reasons,
// creating a mirror view doesn't copy data,
// so we manually do that now.
Kokkos::deep_copy(h_example, example);
// Now we can access the copy of the data on host!
h_example(1) = 4.0;
```

Typical workflow for filling up a View

1. Construct your View
2. Make a host mirror
3. Fill up your host mirror, e.g. with data from a file
4. Deep copy from your host mirror to the View

```
int main() {  
    int num_x_nodes = 100;  
    int num_y_nodes = 100;  
    Kokkos::View<int**> field("field", num_x_nodes, num_y_nodes);  
    Kokkos::View<int**>::HostMirror h_field =  
        Kokkos::create_mirror_view(field);  
    for (int i=0; i<num_x_nodes; i++) {  
        for (int j=0; j<num_y_nodes; j++) {  
            h_field(i, j) = i*j;  
        }  
    }  
    Kokkos::deep_copy(field, h_field);  
}
```

To parallelize for loops, use Kokkos::parallel_for

Here I parallelize the outer for loop of the previous example

- Three arguments for parallel_for
 - The first is a label (just like Views)
 - Second is a range. If it's just a number, the range is 0 to the number, exclusive
 - Third is a **functor**, which is a class or struct with a method that does the body of the desired for loop
- Variables used in the parallel for loop body must be members of the functor
- Note that Views can be accessed inside parallel_fors!
- This is clunky; do we really need to define a class for every single for loop we want to parallelize?

NO. Yay for lambdas!

```
struct Functor{
    // Our functor can have as many member variables as we want,
    // but note that these are all copied to device,
    // so changes do not persist except for Views.
    int num_y_nodes;

    Kokkos::View<int**> field;

    // Functor operators must be marked with the KOKKOS_INLINE_FUNCTION macro
    // and be marked const like this.
    KOKKOS_INLINE_FUNCTION void operator()(const int i) const {
        for (int j=0; j<num_y_nodes; j++) {
            field(i, j) = i*j;
        }
    }
};

int main() {
    Kokkos::ScopeGuard scope_guard;
    int num_x_nodes = 100;
    int num_y_nodes = 100;
    Kokkos::View<int**> field("field", num_x_nodes, num_y_nodes);
    Kokkos::parallel_for(
        "computeField",
        num_x_nodes,
        Functor{num_y_nodes, field}
    );
}
```

Turns out, lambdas ARE functors!

Same functionality, but MUCH simpler!

- So we can write a lambda in place instead of having to define a functor elsewhere
- Kokkos provides a handy KOKKOS_LAMBDA macro that simplifies defining lambdas
- Variables are automatically captured by lambda *by value*
 - Recall that this means variables are copied – effectively, they're copied up to device
 - Since Views are really just pointers to device data, it's okay that they're shallow copied
- Syntax on right attempts to emulate standard for loop syntax as closely as possible
 - It's not exactly the same, but you'll get used to it

```
int main() {
    Kokkos::ScopeGuard scope_guard;
    int num_x_nodes = 100;
    int num_y_nodes = 100;
    Kokkos::View<int**> field("field", num_x_nodes, num_y_nodes);
    Kokkos::parallel_for(
        "computeField",
        num_x_nodes,
        KOKKOS_LAMBDA (const int i)
        {
            for (int j=0; j<num_y_nodes; j++) {
                field(i, j) = i*j;
            }
        });

    Kokkos::View<int**>::HostMirror h_field =
        Kokkos::create_mirror_view(field);
    // Just to compare, here's the serial version again.
    // Note how similar code above is to this.
    // See, parallelism can be easy!
    for (int i=0; i<num_x_nodes; i++) {
        for (int j=0; j<num_y_nodes; j++) {
            h_field(i, j) = i*j;
        }
    }
    Kokkos::deep_copy(field, h_field);
}
```

We can do multidimensional for loops as well

- We just need to use a special **range policy**: Kokkos::MDRangePolicy
 - MD stands for MultiDimensional
- Template argument is a rank, the number of dimensions in range
- First argument is a list of starts for each of the indices
- Second argument is a list of ends for indices, exclusive
- Also must add extra indices as arguments to the lambda

```
int main() {
    Kokkos::ScopeGuard scope_guard;
    int num_x_nodes = 100;
    int num_y_nodes = 100;
    Kokkos::View<int**> field("field", num_x_nodes, num_y_nodes);
    Kokkos::parallel_for(
        "computeField",
        Kokkos::MDRangePolicy<Kokkos::Rank<2>>(
            {0, 0},
            {num_x_nodes, num_y_nodes}
        ),
        KOKKOS_LAMBDA (const int i, const int j)
    {
        field(i, j) = i*j;
    });
}
```

Functions can be used in parallel regions

- But they must be marked with KOKKOS_INLINE_FUNCTION
- CUDA, HIP, and other GPU languages require functions used on device to be marked with special macros
- Kokkos defines KOKKOS_INLINE_FUNCTION to expand into these macros when those backends are enabled
- Otherwise, it expands to nothing

```
KOKKOS_INLINE_FUNCTION
int product(int i, int j) {
    return i*j;
}

int main() {
    Kokkos::ScopeGuard scope_guard;
    int num_x_nodes = 100;
    int num_y_nodes = 100;
    Kokkos::View<int**> field("field", num_x_nodes, num_y_nodes);
    Kokkos::parallel_for(
        "computeField",
        Kokkos::MDRangePolicy<Kokkos::Rank<2>>(
            {0, 0},
            {num_x_nodes, num_y_nodes}
        ),
        KOKKOS_LAMBDA (const int i, const int j)
    {
        field(i, j) = product(i, j);
    });
}
```


Reductions can be parallelized as well

- A reduction is basically just a sum
- Use Kokkos::parallel_reduce
- Add an “update” reference to lambda signature
 - This holds a thread-local temporary sum value
- Result of reduction is saved to last argument in parallel_reduce

```
int main() {  
    Kokkos::View<double*> field(100);  
    Kokkos::deep_copy(field, 1.0);  
    // Want to find the sum of the entries of field  
    double total_sum = 0.0;  
    Kokkos::parallel_reduce(  
        "findSum",  
        field.extent(0), // Size of 0th dimension of field  
        KOKKOS_LAMBDA (const int i, double &update)  
        {  
            update += field(i);  
        }, total_sum);  
    // total_sum now contains the sum of all entries of field  
    return 0;  
}
```

With devices, we can exploit **latency hiding**

- This means host code continues to execute *at the same time* as device code
- Only works if a device backend is enabled
- Fences can be added to synchronize device and host, if needed
 - `deep_copy` has a built-in fence
 - `parallel_reduce` has a built-in fence at the end
- General workflow is to get the `parallel_for` running ASAP, and while it's running, set up data structures that'll use the output

```
int main() {
    Kokkos::ScopeGuard scope_guard;

    Kokkos::View<int*> data("data", 1000000);

    Kokkos::parallel_for(
        "fillData",
        data.extent(0),
        KOKKOS_LAMBDA (const int i)
    {
        data(i) = i;
    });

    // If we're using a device, the following code runs while
    // the parallel_for is still running!
    Kokkos::View<int*>::HostMirror h_data =
        Kokkos::create_mirror_view(data);

    // Use a fence to wait for the device to be done.
    Kokkos::fence();

    // deep_copy also has a built in fence, so no need to fence first.
    Kokkos::deep_copy(h_data, data);

    return 0;
}
```

Must be careful with parallel kernels in class methods

- You should make scope-local copies of member variables you intend to use in parallel kernels
- If you naively use a member variable in a parallel kernel, it'll attempt to copy the ENTIRE class up to device, which is usually not what you want
- Copies of Views are typically shallow, so no need to worry about accidentally copying a million data values

```
struct ParallelExample {  
  
    Kokkos::View<int*> field;  
  
    void computeField() {  
        // This is a shallow copy,  
        // so I'm basically just making a new pointer  
        // to pre-allocated data.  
        Kokkos::View<int*> local_field = field;  
        Kokkos::parallel_for(  
            "computeField",  
            100,  
            KOKKOS_LAMBDA (const int i)  
            {  
                // Note that I'm using local_field  
                // rather than field.  
                // If I used field, the whole class  
                // would be copied up,  
                // and the compiler would likely complain.  
                local_field(i) = i;  
            });  
    }  
};
```

Problem: GPU compilers lack linkers

- That means that code that uses a device function MUST have the definition for that function in the same translation unit
- We can design our file structure to emulate the typical structure, but include device function definitions in headers
- Device function declarations are still in the usual header (marked with KOKKOS_FUNCTION), but definitions are in a separate file, _impl.hpp
- This file is included at the *bottom* of the usual header
- Files that include the header now also get the definitions

Contents of
“DeviceFunctions.hpp”

```
#ifndef DEVICEFUNCTIONS_HPP
#define DEVICEFUNCTIONS_HPP

int hostAdd(int i, int j);

KOKKOS_FUNCTION
int deviceAdd(int i, int j);

#include "DeviceFunctions_impl.hpp"

#endif
```

Contents of
“DeviceFunctions_impl.hpp”

```
#ifndef DEVICEFUNCTIONS_IMPL_HPP
#define DEVICEFUNCTIONS_IMPL_HPP

KOKKOS_INLINE_FUNCTION
int deviceAdd(int i, int j) {
    return i + j;
}

#endif
```

Contents of
“DeviceFunctions.cpp”

```
int hostAdd(int i, int j) {
    return i + j;
}
```

Random number generation is tricky in parallel

- If we give each thread a generator with the same seed, they'll just generate the same numbers
- We have to generate a seed for each thread
- Whenever we need a random number on a thread, we have to grab the seed for that thread, generate the number with that seed, and update the seed so that the next generated number is different
- Ideally the thread-local seeds can all be generated from a single seed, so that the user doesn't need to know the number of threads to specify a seed

Kokkos provides an easier way to do all that

- All in Kokkos_Random.hpp
- First create a random pool
 - The ctor for this generates thread seeds
- Within a parallel kernel, grab a thread-local generator from pool
- Use the generator to generate your random numbers
 - There are several distributions available to draw from
- Must manually release the generator so that it could be grabbed by another thread

```
#include <Kokkos_Core.hpp>
#include <Kokkos_Random.hpp>

int main() {
    Kokkos::ScopeGuard scope_guard;

    // construct a pool with a given seed
    int seed = 1;
    Kokkos::Random_XorShift64_Pool<> rand_pool(seed);

    Kokkos::View<double*> uniform_data("uniform_data", 1000);
    Kokkos::View<double*> normal_data("normal_data", 1000);

    Kokkos::parallel_for(
        "fillData",
        uniform_data.extent(0),
        KOKKOS_LAMBDA (const int i)
    {
        // Grab a generator
        auto rgen = rand_pool.get_state();

        // Generate random numbers uniformly between 2 and 5
        uniform_data(i) = rgen.drand(2.0, 5.0);

        // Generate random numbers distributed Gaussian
        // with mean 1 and standard deviation 2
        normal_data(i) = rgen.normal(1.0, 2.0);

        // Release generator
        rand_pool.free_state(rgen);
    });

    return 0;
}
```

ScatterViews help to parallelize histograms

- In the Experimental namespace
- Charge density is basically a histogram
- ScatterView is created based on a standard View “target”
- Behavior depends on DefaultExecutionSpace
 - With devices, ScatterView is just a shallow copy of target with **atomic access** – only one thread is allowed to write at a time
 - Without devices, ScatterView is a deep copy of the target per thread, so each thread has its own array. These need to be combined later on
- Generally, contribute ASAP after histogramming, and reset right before the next histogram

```
#include <Kokkos_ScatterView.hpp>

int main() {
    Kokkos::ScopeGuard scope_guard;
    int num_nodes = 100;

    Kokkos::View<double*> charge("charge", num_nodes);
    Kokkos::Experimental::ScatterView<double*> scatter_charge(charge);

    // Add into the ScatterView
    Kokkos::parallel_for(
        "addIntoScatterView",
        num_nodes,
        KOKKOS_LAMBDA (const int i)
    {
        // Need to get an "access" object to ScatterView
        auto access = scatter_charge.access();

        // Now add into access as if it were a regular array!
        access(i) += 10.0 * i;
    });

    // Need to contribute to the target View to consolidate histogram
    Kokkos::Experimental::contribute(charge, scatter_charge);

    // We can zero out all the arrays with the reset method
    scatter_charge.reset();
}
```

Class methods don't work easily on device

- When copying class to device, all the member variables get copied as desired
- But when you mark a method `KOKKOS_INLINE_FUNCTION`, both a host and device version are compiled
- Only the host methods get copied to device if you naively copy a class
- When attempting to call method, function pointer is invalid – CUDA will emit a *very* subtle error!

Example of naive class copy that fails

```
class InvalidMethod {
public:
    double to_add = 3.0;

    KOKKOS_INLINE_FUNCTION
    double increment(double to_increment) const {
        return to_add + to_increment;
    }
};

int main() {
    Kokkos::ScopeGuard scope_guard;

    InvalidMethod invalid_method;
    Kokkos::parallel_for(
        "attemptInvalidMethod",
        100,
        KOKKOS_LAMBDA (const int i)
        {
            // This is okay
            double to_add = invalid_method.to_add;

            // Invalid!
            double incremented = invalid_method.increment(i);
        });

    return 0;
}
```


Placement new lets us sidestep this issue

1. Pre-allocate enough space for the class and store a pointer to it
 2. Create a new class instance at that pointer
 3. Copy only member variables to new class
- Since the class was created on device, it has the correct device function pointers
 - All the member variables are preserved as desired

```
class ValidMethod {
public:
    double to_add = 3.0;

    KOKKOS_INLINE_FUNCTION
    double increment(double to_increment) const {
        return to_add + to_increment;
    }
};

int main() {
    Kokkos::ScopeGuard scope_guard;

    ValidMethod valid_method;

    // Step 1: preallocate space on device and store pointer
    ValidMethod *d_valid_method = static_cast<ValidMethod*>{
        Kokkos::kokkos_malloc(sizeof(ValidMethod))
    };
    // Step 2 and 3: create new class instance at pointer,
    // copy members to new instance.
    // We enter a 1-long parallel_for just to get on device.
    Kokkos::parallel_for(
        "copyClass",
        1,
        KOKKOS_LAMBDA (const char)
        {
            // This line handles creation at pointer and copying,
            // using a copy constructor (which I haven't talked about).
            new (d_valid_method) ValidMethod(valid_method);
        });

    // Now we can use d_valid_method as a pointer to the copy of the class.
    Kokkos::parallel_for(
        "attemptValidMethod",
        100,
        KOKKOS_LAMBDA (const int i)
        {
            // Valid!
            double incremented = d_valid_method->increment(i);
        });
    Kokkos::fence();

    // Must manually free data at pointer.
    // In principle, this can be done automatically with
    // a smart pointer.
    Kokkos::kokkos_free(d_valid_method);

    return 0;
}
```

SYCL does not natively support virtual functions on device

- Specifically it doesn't support **indirect function calls** – basically any call where the function could change
- We can get around this with type enumeration
- Basically we build our own runtime derived-class type determination into the classes themselves
- Base class has a public member variable that determines the type
- Derived classes store the right type in that variable
- Non-member function can read that variable, cast the base class to the appropriate derived class, and call the “virtual” method

An example of
type enumeration
on host

```
#include <iostream>

enum class type_t {
    CHILD_A,
    CHILD_B,
};

struct Base {
    type_t type;

    Base(type_t type)
        : type(type) {}
};

struct ChildA : public Base {
    ChildA()
        : Base(type_t::CHILD_A) {}

    void pseudo_virtual() {
        std::cout << "A calling." << std::endl;
    }
};

struct ChildB : public Base {
    ChildB()
        : Base(type_t::CHILD_B) {}

    void pseudo_virtual() {
        std::cout << "B calling." << std::endl;
    }
};

void call_pv(Base* base) {
    if (base->type == type_t::CHILD_A) {
        static_cast<ChildA*>(base)->pseudo_virtual();
    } else {
        static_cast<ChildB*>(base)->pseudo_virtual();
    }
}
```

Gotta add some fluff to make it work on device

- Just preceding every function that'll be called on device with KOKKOS_INLINE_FUNCTION
- Combine this with placement new to use pseudo-virtual functions on device!

```
#include <iostream>
#include <Kokkos_Core.hpp>

enum class type_t {
    CHILD_A,
    CHILD_B,
};

You, 44 seconds ago | 2 authors (You and others)
struct Base {
    type_t type;

    Base(type_t type)
    : type(type) {}
};

You, 44 seconds ago | 1 author (You)
struct ChildA : public Base {
    ChildA()
    : Base(type_t::CHILD_A) {}

    KOKKOS_INLINE_FUNCTION
    void pseudo_virtual() {
        std::cout << "A calling." << std::endl;
    }
};

You, 44 seconds ago | 1 author (You)
struct ChildB : public Base {
    ChildB()
    : Base(type_t::CHILD_B) {}

    KOKKOS_INLINE_FUNCTION
    void pseudo_virtual() {
        std::cout << "B calling." << std::endl;
    }
};

KOKKOS_INLINE_FUNCTION
void call_pv(Base* base) {
    if (base->type == type_t::CHILD_A) {
        static_cast<ChildA*>(base)->pseudo_virtual();
    } else {
        static_cast<ChildB*>(base)->pseudo_virtual();
    }
}
```

Exercise: write a 1D particle pusher!

For now, we'll just use a constant force. If you have a working hPIC2 installation, you can just delete the contents of main.cpp and write this here so that you have easy access to Kokkos libraries.

1. Use previous lessons to write a main function. You can just hardcode the parameters (number of particles, particle weight, number of grid elements, grid size, number of timesteps, timestep size, constant acceleration) and make it a main function with no inputs.
2. Use initialize/finalize or a ScopeGuard to set up a Kokkos instance.
3. Make 2 `View<double*>`, one for particle positions and one for particle velocities. Make an RNG pool from a hardcoded seed, then enter a `parallel_for` to randomly generate particle positions uniformly throughout domain and particle velocities from a Gaussian distribution.
4. Make a `View<double*>` for the charge accumulated by particles at each node, and make a `ScatterView` targeting this.
5. Write an outer for loop over timesteps
6. At each timestep...
 - a. Reset the charge `ScatterView`
 - b. Do a `parallel_for` over all particles
 - i. Update velocity, $v += a * dt$
 - ii. Update position, $x += v * dt$
 - iii. If a particle leaves the grid, just wrap it back around to the other side Pacman-style (periodic boundaries)
 - iv. Find the two nodes adjacent to the element containing the particle, linearly interpolate particle weight to each node, get access to charge `ScatterView`, and contribute charge to those nodes in the `ScatterView`
 - c. contribute charge from `ScatterView` to target `View`
7. Compile with multiple Kokkos backends to compare speeds between, e.g., OpenMP and serial